

Getting Started
Frequently Asked Questions
Part 1
Introduction to Part 1
Definitions and basic concepts
Running and stopping containers
In-depth dive into images
Defining start conditions for the container
Interacting with the container via volumes and ports
Utilizing tools from the Registry
Summary
Part 2
Part 3

In-depth dive into images

Images are the basic building blocks for containers and other images. When you "containerize" an application you work towards creating the image.

By learning what images are and how to create them you are ready to start utilizing containers in your own projects.

Where do the images come from?

When running a command such as `docker run hello-world`, Docker will automatically search [Docker Hub](#) for the image if it is not found locally.

This means that we can pull and run any public image from Docker's servers. For example, if we wanted to start an instance of the PostgreSQL database, we could just run `docker run postgres`, which would pull and run https://hub.docker.com/_/postgres.

We can search for images in the Docker Hub with `docker search`. Try running `docker search hello-world`.

The search finds plenty of results, and prints each image's name, short description, amount of stars, and "official" and "automated" statuses.

```
$ docker search hello-world
NAME                    DESCRIPTION           STARS     OFFICIAL   AUTOMATED
hello-world              Hello World!...       1988     [OK]
kitematic/hello-world-nginx A light-weight... 153
tutum/hello-world        Image to test...    90          [OK]
...
```

Let's examine the list.

The first result, `hello-world`, is an official image. [Official images](#) are curated and reviewed by Docker, Inc. and are usually actively maintained by the authors. They are built from repositories in the [docker-library](#).

When browsing the CLI's search results, you can recognize an official image from the "[OK]" in the "OFFICIAL" column and also from the fact that the image's name has no prefix (aka organization/user). When browsing Docker Hub, the page will show "Docker Official Images" as the repository, instead of a user or organization. For example, see the [Docker Hub page of the hello-world image](#).

The third result, `tutum/hello-world`, is marked as "automated". This means that the image is automatically built from the source repository. Its [Docker Hub page](#) shows its previous "Builds" and a link to the image's "Source Repository" (in this case, to GitHub) from which Docker Hub builds the image.

There are also other Docker registries competing with Docker Hub, such as [Quay](#). By default, `docker search` will only search from Docker Hub, but to a search different registry, you can add the registry address before the search term, for example, `docker search quay.io/hello`. Alternatively, you can use the registry's web pages to search for images. Take a look at the page of the [nordstrom/hello-world image on Quay](#). The page shows the command to use to pull the image, which reveals that we can also pull images from hosts other than Docker Hub:

```
docker pull quay.io/nordstrom/hello-world
```

So, if the host's name (here: `quay.io`) is omitted, it will pull from Docker Hub by default.

NOTE: Trying the above command may fail giving manifest errors as the default tag `latest` is not present in `quay.io/nordstrom/hello-world` image. Specifying a correct tag for a image will pull the image without any errors, for ex. `docker pull quay.io/nordstrom/hello-world:2.0`

A detailed look into an image

Let's go back to a more relevant image than 'hello-world', the Ubuntu image, one of the most common Docker images to use as a base for your own image.

Let's pull Ubuntu and look at the first lines:

```
$ docker pull ubuntu
Using default tag: latest
latest: Pulling from library/ubuntu
```

Since we didn't specify a tag, Docker defaulted to `latest`, which is usually the latest image built and pushed to the registry. However, in this case, the repository's [README](#) says that the `ubuntu:latest` tag points to the "latest LTS" instead since that's the version recommended for general use.

Images can be tagged to save different versions of the same image. You define an image's tag by adding `:<tag>` after the image's name.

Ubuntu's [Docker Hub page](#) reveals that there's a tag named `22.04` which promises us that the image is based on Ubuntu 22.04. Let's pull that as well:

```
$ docker pull ubuntu:22.04
22.04: Pulling from library/ubuntu
c2ca09a1934b: Downloading [=====] 34.25MB/38.64MB
d6c3619d2153: Download complete
0ef07335a04: Download complete
6b1bb01b3a3b: Download complete
43a98c187399: Download complete
```

Images are composed of different layers that are downloaded in parallel to speed up the download. Images being made of layers also have other aspects and we will talk about them in part 3.

We can also tag images locally for convenience, for example, `docker tag ubuntu:22.04 ubuntu:jammy_jellyfish` creates the tag `ubuntu:jammy_jellyfish` which refers to `ubuntu:22.04`.

Tagging is also a way to "rename" images. Run `docker tag ubuntu:22.04 fav_distro:jammy_jellyfish` and check `docker image ls` to see what effects the command had.

To summarize, an image name may consist of 3 parts plus a tag. Usually like the following: `registry/organisation/image:tag`. But may be as short as `ubuntu`, then the registry will default to Docker hub, organisation to `library` and tag to `latest`. The organisation may also be a user, but calling it an organisation may be more clear.

Exercises 1.5 - 1.6

EXERCISE 1.5: SIZES OF IMAGES

In the Exercise 1.3 we used `devopsdockeruh/simple-web-service:ubuntu`.

Here is the same application but instead of Ubuntu is using Alpine Linux: `devopsdockeruh/simple-web-service:alpine`.

Pull both images and compare the image sizes. Go inside the Alpine container and make sure the secret message functionality is the same. Alpine version doesn't have `bash` but it has `sh`, a more bare-bones shell.

EXERCISE 1.6: HELLO DOCKER HUB

Run `docker run -it devopsdockeruh/pull_exercise`.

The command will wait for your input.

Navigate through the Docker hub to find the docs and Dockerfile that was used to create the image.

Read the Dockerfile and/or docs to learn what input will get the application to answer a "secret message".

Where do the images come from?

A detailed look into an image

Exercises 1.5 - 1.6

Building images

Exercises 1.7 - 1.8

Submit the secret message and command(s) given to get it as your answer.

Building images

Finally, we get to build our own images and get to talk about [Dockerfile](#) and why it's so great.

Dockerfile is simply a file that contains the build instructions for an image. You define what should be included in the image with different instructions. We'll learn about the best practices here by creating one.

Let's take a most simple application and containerize it first. Here is a script called "hello.sh"

hello.sh

```
#!/bin/sh
echo "Hello, docker!"
```

First, we will test that it even works. Create the file, add execution permissions and run it:

```
$ chmod +x hello.sh
$ ./hello.sh
Hello, docker!
```

- If you're using Windows you can skip these two and add chmod +x hello.sh to the Dockerfile.

And now to create an image from it. We'll have to create the [Dockerfile](#) that declares all of the required dependencies. At least it depends on something that can run shell scripts. We will choose [Alpine](#), a small Linux distribution that is often used to create small images.

Even though we're using Alpine here, you can use Ubuntu during exercises. Ubuntu images by default contain more tools to debug what is wrong when something doesn't work. In part 3 we will talk more about why small images are important.

We will choose exactly which version of a given image we want to use. This guarantees that we don't accidentally update through a breaking change, and we know which images need updating when there are known security vulnerabilities in old images.

Now create a file and name it "Dockerfile" and put the following instructions inside it:

Dockerfile

```
# Start from the alpine image that is smaller but no fancy tools
FROM alpine:3.19

# Use /usr/src/app as our workdir. The following instructions will be executed in this location.
WORKDIR /usr/src/app

# Copy the hello.sh file from this directory to /usr/src/app/ creating /usr/src/app/hello.sh
COPY hello.sh .

# Alternatively, if we skipped chmod earlier, we can add execution permissions during the build.
# RUN chmod +x hello.sh

# When running docker run the command will be ./hello.sh
CMD ./hello.sh
```

Great! We can use the command [docker build](#) to turn the Dockerfile to an image.

By default [docker build](#) will look for a file named Dockerfile. Now we can run [docker build](#) with instructions where to build (.,) and give it a name (-t <name>):

```
$ docker build -t hello-docker .
=> [internal] load build definition from Dockerfile
=> transferring Dockerfile: 478B
=> [internal] load metadata for docker.io/library/alpine:3.19
=> [auth] library/alpine:pull token for registry-1.docker.io
=> [internal] load .dockerignore
=> [internal] transfer context: 2B
=> [1/3] FROM docker.io/library/alpine:3.19@sha256:c5b1261d6d3e43071626931fc004f70149baeba2c8ec672bd4f
=> [internal] load build context
=> [2/3] WORKDIR /usr/src/app
=> [3/3] COPY hello.sh .
=> exporting image
=> exporting layers
=> writing image sha256:5f8f5d7445f34b0bcfaaa4d685a068cdcc1ed79e650683373a228c79ea69c8
=> naming to docker.io/library/hello-docker
```

Let us ensure that the image exists:

```
$ docker image ls
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
hello-docker        latest   5f8f5d7445f3   4 minutes ago  7.73MB
```

⚠ PERMISSION DENIED

If you're now getting "/bin/sh: ./hello.sh: Permission denied" it's because the [chmod +x hello.sh](#) was skipped earlier. You can simply uncomment the RUN instruction between COPY and CMD instructions

⚠ NOT FOUND

If you're now getting "/bin/sh: ./hello.sh: not found" and you're using Windows it might be because by default Windows uses CRLF as line ending. Unix, in our case Alpine, uses just LF which makes the copying of our hello.sh invalid bash script in the build phase. To overcome this error change the line endings to LF before running [docker build](#).

Now executing the application is as simple as running [docker run hello-docker](#). Try it!

During the build we see from the output that there are three steps: [1/3], [2/3] and [3/3]. The steps here represent layers of the image so that each step is a new layer on top of the base image (alpine:3.19 in our case).

Layers have multiple functions. We often try to limit the number of layers to save on storage space but layers can work as a cache during build time. If we just edit the last line of Dockerfile the build command can start from the previous layer and skip straight to the section that has changed. COPY automatically detects changes in the files, so if we change the hello.sh it'll run from step 3/3, skipping 1 and 2. This can be used to create faster build pipelines. We'll talk more about optimization in part 3.

It is also possible to manually create new layers on top of a image. Let us now create a new file called [additional.txt](#) and copy it inside a container.

We'll need two terminals, that shall be called 1 and 2 in the following listings. Let us start by running the image:

```
# do this in terminal 1
$ docker run -it hello-docker sh
/usr/src/app #
```

Now we're inside of the container. We replaced the CMD we defined earlier with [sh](#) and used -i and -t to start the container so that we can interact with it.

In the second terminal we will copy the file inside the container:

```
# do this in terminal 2
$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS           NAMES
9c0695e3e85        hello-docker        "sh"              4 minutes ago    Up 4 minutes          zen_rosalind

$ touch additional.txt
$ docker cp ./additional.txt zen_rosalind:/usr/src/app/
```

The file is created with command [touch](#) right before copying it in.

Let us ensure that the file is copied inside the container:

```
# do this in terminal 1
/usr/src/app # ls
additional.txt  hello.sh
```

Great! Now we've made a change to the container. We can use command `docker diff` to check what has changed

```
# do this in terminal 2
$ docker diff zen_rosalind
C /usr
C /usr/src
C /usr/src/app
A /usr/src/app/additional.txt
C /root
A /root/.ash_history
```

The character in front of the file name indicates the type of the change in the container's filesystem: A = added, D = deleted, C = changed. The `additional.txt` was created and our `.ash_history`.

Next we will save the changes as a new image with the command `docker commit`:

```
# do this in terminal 2
$ docker commit zen_rosalind hello-docker-additional
sha256:2f63baa355ce5970b789fe000b92717f25d91172ed71620be784315fc4fd
$ docker image ls
REPOSITORY          TAG      IMAGE ID      CREATED        SIZE
hello-docker-additional    latest   2f63baa355ce  3 seconds ago  7.73MB
hello-docker           latest   444f21cf7bd5  31 minutes ago  7.73MB
```

Technically the command `docker commit` added a new layer on top of the image `hello-docker`, and the resulting image was given the name `hello-docker-additional`.

We will actually not use the command `docker commit` again during this course. This is because defining the changes to the Dockerfile is much more sustainable method of managing changes. No magic actions or scripts, just a Dockerfile that can be version controlled.

Let's do just that and create `hello-docker` with v2 tag that includes the file `additional.txt`. The new file can be added with a `RUN` instruction:

Dockerfile

```
# Start from the alpine image
FROM alpine:3.19

# Use /usr/src/app as our workdir. The following instructions will be executed in this location.
WORKDIR /usr/src/app

# Copy the hello.sh file from this location to /usr/src/app/ creating /usr/src/app/hello.sh.
COPY hello.sh .

# Execute a command with '/bin/sh -c' prefix.
RUN touch additional.txt

# When running Docker run the command will be ./hello.sh
CMD ./hello.sh
```

Now we used the `RUN` instruction to execute the command `touch additional.txt` which creates a file inside the resulting image. Pretty much anything that can be executed in the container based on the created image, can be instructed to be run with the `RUN` instruction during the build of a Dockerfile.

Build now the Dockerfile with `docker build . -t hello-docker:v2` and we are done! Let's compare the output of `ls`:

```
$ docker run hello-docker-additional ls
additional.txt
hello.sh

$ docker run hello-docker:v2 ls
additional.txt
hello.sh
```

Now we know that all instructions in a Dockerfile **except** `CMD` (and one other that we will learn about soon) are executed during build time. `CMD` is executed when we call `docker run`, unless we overwrite it.

Exercises 1.7 - 1.8

① EXERCISE 1.7: IMAGE FOR SCRIPT

We can improve our previous solutions now that we know how to create and build a Dockerfile.

Let us now get back to [Exercise 1.4](#).

Create a new file `script.sh` on your local machine with the following contents:

```
while true
do
  echo "Input website:"
  read website; echo "Searching.."
  sleep 1; curl http://$website
done
```

Create a Dockerfile for a new image that starts from `ubuntu:22.04` and add instructions to install `curl` into that image. Then add instructions to copy the script file into that image and finally set it to run on container start using `CMD`.

After you have filled the Dockerfile, build the image with the name "curler".

- If you are getting permission denied, use `chmod` to give permission to run the script.

The following should now work:

```
$ docker run -it curler
Input website:
helsinki.fi
Searching..
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<html><head>
<title>301 Moved Permanently</title>
</head><body>
<h1>Moved Permanently</h1>
<p>The document has moved <a href="https://www.helsinki.fi/">here</a>.</p>
</body></html>
```

Remember that `RUN` can be used to execute commands while building the image!

Submit the Dockerfile.

① EXERCISE 1.8: TWO LINE DOCKERFILE

By default our `devopsdockerrh/simple-web-service:alpine` doesn't have a `CMD`. Instead, it uses `ENTRYPOINT` to declare which application is run.

We'll talk more about `ENTRYPOINT` in the next section, but you already know that the last argument in `docker run` can be used to give a command or an argument.

As you might've noticed it doesn't start the web service even though the name is "simple-web-service". A suitable argument is needed to start the server!

Try `docker run devopsdockerrh/simple-web-service:alpine hello`. The application reads the argument "hello" but will inform that hello isn't accepted.

In this exercise create a Dockerfile and use `FROM` and `CMD` to create a brand new image that automatically runs `server`.

The Docker documentation `CMD` says a bit indirectly that if a image has `ENTRYPOINT` defined, `CMD` is used to define it the default arguments.

Tag the new image as "web-server"

Return the Dockerfile and the command you used to run the container.

running the built "web-server" image should look like this:

```
$ docker run web-server
[GIN-debug] [WARNING] Creating an Engine instance with the Logger and Recovery middleware already attached. This will result in a memory leak in production unless you are using gin.ReleaseMode()
[GIN-debug] [WARNING] Running in "debug" mode. Switch to "release" mode in production.
- using env: export GIN_MODE=release
- using code: gin.SetMode(gin.ReleaseMode)

[GIN-debug] GET    /*path                           --> server.Start.func1 (3 handlers)
[GIN-debug] Listening and serving HTTP on :8080
```

- We don't have any method of accessing the web service yet. As such confirming that the console output is the same will suffice.
- The exercise title may be a useful hint here.

[Edit this page](#)

Previous

[« Running and stopping containers](#)

Next

[Defining start conditions for the container »](#)

Help

[Discord](#)

[Report an issue](#)

More

[About](#)

[GitHub](#)

In collaboration

[University of Helsinki](#)

[Epicode](#)

[Unity](#)